

Lab 06 – Two-sample t -test

ENVX1002 Handbook

Semester 1, 2025

Learning outcomes

- Learn to use R to calculate a 2-sample t -test
 - independent samples with constant variance
 - independent samples with unequal variance
 - paired samples
 - data transformations
- Apply the steps for hypothesis testing from lectures
- Learn how to interpret statistical output

Before you begin

Create your Quarto document and save it as Lab-06.Rmd or similar. The following data files are required:

- 1) Barley.csv
- 2) Plant_growth.csv
- 3) Turbidity.csv

The following external packages are used in this lab. Install them if you have not done so already.

```
install.packages(c("tidyverse", "car"),  
  repo = "https://cloud.r-project.org")
```

Finally, try to complete today's lab exercises in pairs and try out pair programming, where one person writes the code and the other person reviews each line as it is written. You can swap roles every 10 minutes or so. This is a great way to learn from each other and to improve your coding skills.

Exercise 1: barley (walk-through)

Background

An experiment was designed to compare two varieties of spring barley. Thirty four plots were used, seventeen being randomly allocated to variety A and seventeen to variety B. Unfortunately

five plots were destroyed. The yields (t/ha) from the remaining plots were as they appear in the file `Barley.csv`.

Instructions

First, quickly explore the data; then, utilise the **HATPC** process and test the hypothesis that the two varieties give equal yields, assuming that the samples are independent.

HATPC:

- Hypothesis
- Assumptions
- Test (statistic)
- P-value
- Conclusion

i Level of significance

The level of significance is usually set at 0.05. This value is generally accepted in the scientific community and is also linked to Type 2 errors, where choosing a lower significance increases the likelihood of failing to reject the null hypothesis when it is false.

Data exploration

First we load the data and inspect its structure to see if it needs to be cleaned or transformed. The `glimpse()` function is a tidy version of `str()` that provides a quick overview of the data that focuses on the variables, ignoring data attributes.

Try to compare `str()` and `glimpse()` to see what the differences are.

```
barley <- readr::read_csv("data/Barley.csv") # packagename:: before a function  
lets you access a function without having to load the whole library first
```

```
Rows: 29 Columns: 2  
— Column specification  
Delimiter: ","  
chr (1): Variety  
dbl (1): Yield  
  
i Use `spec()` to retrieve the full column specification for this data.  
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
dplyr::glimpse(barley)
```


```
Rows: 29
Columns: 2
$ Variety <chr> "A", "A", "A", "A", "A", "A", "A", "A", "A", "A", "A", "A", "A", "A...
$ Yield <dbl> 4.6, 4.3, 3.8, 3.4, 3.9, 3.9, 3.9, 4.4, 3.6, 3.6, 4.7, 3.9, 3.9...
```

The Variety column is a factor with two levels, A and B, but it is defined as a character. We can convert it to a factor using the `mutate()` function from the `dplyr` package, but it is not necessary for the *t*-test since R will automatically convert it to a factor.

```
library(tidyverse)
barley <- mutate(barley, Variety = as.factor(Variety))
```

Quickly preview the data as a plot to see if there are any trends or unusual observations.

```
barley %>%
  ggplot(aes(x = Variety, y = Yield)) +
  geom_boxplot()
```



A trained eye will anticipate that the data may not meet the assumption of equal variance; however, we will test this assumption later. Otherwise, there appear to be no unusual observations in the data.

Hypothesis

What are the null and alternative hypotheses? We can use the following notation:

$$H_0 : \mu_A = \mu_B$$

$$H_1 : \mu_A \neq \mu_B$$

where μ_A and μ_B are the population means of varieties A and B, respectively.

It is **important** that when using mathematical symbols to denote the null and alternative hypotheses, you should **always** define what the symbols mean. Otherwise, the reader may not understand what you are referring to.

The equations above are written in LaTeX, a typesetting system that is commonly used in scientific writing. You can learn more about LaTeX [here](#). The raw syntax used to write the equations are shown below:

```
$$H_0: \mu_A = \mu_B$$
$$H_1: \mu_A \neq \mu_B$$
```

Why do we always define the null and alternative hypotheses? In complex research projects or when working in a team, it is important to ensure that everyone is on the same page. By defining the hypotheses, you can avoid misunderstandings and ensure that everyone is working towards the same goal as the mathematical notation is clear and unambiguous.

Assumptions

Normality

There are many ways to check for normality. Here we will use the QQ-plot. Use of ggplot2 is preferred (as a means of practice) but since we are just exploring data, base R functions are not a problem to use.

Using ggplot2

```
ggplot(barley, aes(sample = Yield)) +
  stat_qq() +
  stat_qq_line() +
  facet_wrap(~Variety) #facet_wrap ensures there are separate plots for each
  variety rather than one plot with all the data in Yield
```



Using base R

```
par(mfrow = c(1, 2))
qqnorm(barley$Yield[barley$Variety == "A"], main = "Variety A") # square brackets
to subset the data by variety
qqline(barley$Yield[barley$Variety == "A"])
qqnorm(barley$Yield[barley$Variety == "B"], main = "Variety B")
qqline(barley$Yield[barley$Variety == "B"])
```



Question: Do the plots indicate the data are normally distributed?

Answer: Yes, the data appear to be normally distributed as the QQ-plot shows that the data points are close to the line.

Homogeneity of variance

From the boxplot, we can see that there is some indication that the variances are not equal. We can test this assumption using Bartlett's test or Levene's test; here we will just use Bartlett's test.

```
bartlett.test(Yield ~ Variety, data = barley)
```

Bartlett test of homogeneity of variances

data: Yield by Variety

Bartlett's K-squared = 14.616, df = 1, p-value = 0.0001318

Question: Does the Bartlett's test indicate the two groups have equal variances? What effect will that have on the analysis?

Answer: The two groups have unequal variance (Bartlett's test: $X^2 = 14.6, p < 0.01$). This means that we will need to use the Welch's t -test, which does not assume equal variances.

Test statistic

We can now calculate the test statistic using the `t.test()` function in R. Since the variances are unequal, we do not have to specify the `var.equal` argument – the default test for `t.test()` is the Welch's *t*-test which does not assume equal variances.

```
t.test(Yield ~ Variety, data = barley)
```

```
Welch Two Sample t-test

data: Yield by Variety
t = -4.9994, df = 19.441, p-value = 7.458e-05
alternative hypothesis: true difference in means between group A and group B is
not equal to 0
95 percent confidence interval:
 -0.9293569 -0.3814274
sample estimates:
mean in group A mean in group B
    4.052941      4.708333
```

P-value

Since the p-value is < 0.05 , we can reject the null hypothesis that the mean yield of both varieties is equal.

Conclusion

The conclusion needs to be brought into the context of the study. In a scientific report or paper, you would write something like this:

The mean yield of barley variety A was significantly different from that of variety B ($t = -5.0$, $df = 19.4$, $p < 0.01$).

Exercise 2: plant growth

Background

In a test of a particular treatment aimed at inducing growth, 20 plants were grouped into ten pairs so that the two members of each pair were as similar as possible. One plant of each pair was chosen randomly and treated; the other was left as a control. The increases in height (in centimetres) of plants over a two-week period are given in the file `Two week plant heights`. We wish to compare whether the treatment is actually inducing improved growth, as compared to the control.

Instructions

Here, we have two samples, and the samples are paired as it is a before-after experiment. So we'd like to conduct a paired *t*-test.

For paired t -tests the analysis is performed as a 1-sample t -test on the difference between each pair so the only assumption is the normality assumption.

Copy the structure below and perform your analysis in your document.

```
## Exercise 2: plant growth
### Data exploration
### Hypothesis
### Assumptions
#### Normality
#### Homogeneity of variance
### Test statistic
### P-value
### Conclusion
```

Note that the data is *not* tidy. The code below will convert the data to the long format and assign it to `tidy_plant`.

```
plant_growth <- readr::read_csv("data/Plant_growth.csv")
```

```
Rows: 10 Columns: 3
—
Column specification
Delimiter: ","
dbl (3): Pair, Treated, Control

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
tidy_plant <- plant_growth %>%
  pivot_longer(cols = c(Treated, Control), names_to = "Treatment", values_to = "Height")
```

You may also need to perform a Shapiro-Wilk test to check for normality. To do this for each group, you can use the `tapply()` function.

```
tapply(tidy_plant$Height, tidy_plant$Treatment, shapiro.test)
```

```
$Control

Shapiro-Wilk normality test

data:  X[[i]]
W = 0.89588, p-value = 0.1973
```



```
$Treated


Shapiro-Wilk normality test

data:  X[[i]]
W = 0.93139, p-value = 0.4617
```

Answer 2

Data exploration

```
#Boxplot
ggplot(tidy_plant, aes(x = Treatment, y = Height)) +
  geom_boxplot()
```



The boxplot shows no unusual observations.

Hypothesis

$$H_0 : \mu_{\text{treated}} = \mu_{\text{control}}$$


$$H_1 : \mu_{\text{treated}} \neq \mu_{\text{control}}$$

where μ_{treated} and μ_{control} are the mean increases in height of the treated and control plants, respectively.

Assumptions

Normality

```
#qqplots  
ggplot(tidy_plant, aes(sample = Height)) +  
  stat_qq() +  
  stat_qq_line() +  
  facet_wrap(~Treatment)
```



The QQ-plots show that the data are normally distributed, but there are some outliers worth checking. A formal test for normality will give us a better idea of whether the assumption is met.

```
# perform test on groups  
tapply(tidy_plant$Height, tidy_plant$Treatment, shapiro.test)
```

\$Control

Shapiro-Wilk normality test

data: X[[i]]

W = 0.89588, p-value = 0.1973

\$Treated

Shapiro-Wilk normality test

```
data: X[[i]]  
W = 0.93139, p-value = 0.4617
```

The Shapiro-Wilk test indicates that the data are normally distributed.

Homogeneity of variance

There is **no need** to test for homogeneity of variance in a paired *t*-test!

Test statistic

```
t.test(tidy_plant$Height[tidy_plant$Treatment == "Treated"],  
       tidy_plant$Height[tidy_plant$Treatment == "Control"],  
       paired = TRUE  
)
```

Paired t-test

```
data:      tidy_plant$Height[tidy_plant$Treatment == "Treated"] and  
tidy_plant$Height[tidy_plant$Treatment == "Control"]  
t = 2.5468, df = 9, p-value = 0.03136  
alternative hypothesis: true mean difference is not equal to 0  
95 percent confidence interval:  
 0.1900113 3.2099887  
sample estimates:  
mean difference  
      1.7
```

P-value

The p-value is 0.03 which is less than 0.05, so we reject the null hypothesis that the mean increase in height of the treated and control plants is equal.

Conclusion

The mean increase in height of the treated plants is significantly different from that of the control plants ($t = 2.5$, $df = 9$, $p = 0.03$).

Exercise 3: turbidity

Background

A new filtering process was installed at a dam which provided drinking water for a nearby town. To check on its success, a number of water samples were taken at random times and locations in the weeks before and after the process was installed. The following are the turbidity values (units = NTU) of the water samples.

Instructions

Now we consider further examples of a two-sample t -test, but where the assumption of equal variance and normality may not be met for the raw data. Sometimes after applying a data transformation the analysis can proceed assuming equal variances – but always check after a data transformation.

The data can be read with the code below:

```
turbidity <- read_csv("data/Turbidity.csv")
```

```
Rows: 19 Columns: 2
—                Column                specification
-----
Delimiter: ",",
chr (1): Time
dbl (1): Turbidity

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

For data transformation, you may need to create a new variable in your dataset to store the transformed data. For example, to create a new variable `TurbLog10` that stores the \log_{10} transformed turbidity values, you can use the following code:

```
turbidity$TurbLog10 <- log10(turbidity$Turbidity)
```

To interpret the results for your conclusions, you may need to back-transform the mean and/or confidence interval values. To back transform \log_{10} data you use:

$$10^{\text{mean or CI}}$$

To back-transform natural log, \log_e , you use:

$$e^{\text{mean or CI}}$$

Answer 3

Data exploration


```
#Read in data
turbidity <- read_csv("data/Turbidity.csv")
```

```
Rows: 19 Columns: 2
—                Column                specification
-----
Delimiter: ",",
```

```
chr (1): Time
dbl (1): Turbidity
```

```
i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
#Make boxplot
ggplot(turbidity, aes(x = Time, y = Turbidity)) +
  geom_boxplot()
```



The boxplot shows skewing of the data, and outliers, a strong indication that the data will have issues meeting the assumption of normality. It is also likely that the variances are not equal as the boxplot shows different spreads.

Hypothesis

$$H_0 : \mu_{\text{before}} = \mu_{\text{after}}$$


$$H_1 : \mu_{\text{before}} \neq \mu_{\text{after}}$$

where μ_{before} and μ_{after} are the mean turbidity values before and after the new filtering process was installed.

Assumptions

Normality: part 1

```
#qqplots
ggplot(turbidity, aes(sample = Turbidity)) +
  stat_qq() +
  stat_qq_line() +
  facet_wrap(~Time)
```



The QQ-plots show that the data are not normally distributed. Transformation of the data is required – there is no need to look at the equal variances just yet.


Transform data

We will use the natural log transformation to normalise the data.

```
turbidity$TurbLog10 <- log10(turbidity$Turbidity)
```

Normality: part 2

```
ggplot(turbidity, aes(sample = TurbLog10)) +
  stat_qq() +
  stat_qq_line() +
  facet_wrap(~Time)
```



The QQ-plots show that the data look better. We can now perform a Shapiro-Wilk test to confirm this.

```
tapply(turbidity$TurbLog10, turbidity$Time, shapiro.test)
```

```
$After
      Shapiro-Wilk normality test

data:  X[[i]]
W = 0.96995, p-value = 0.8981

$Before
      Shapiro-Wilk normality test

data:  X[[i]]
W = 0.97566, p-value = 0.9602
```

The Shapiro-Wilk test indicates that the data are normally distributed.

Homogeneity of variance

We check the assumption of equal variance after the transformation.

```
bartlett.test(TurbLog10 ~ Time, data = turbidity)
```

Bartlett test of homogeneity of variances

```
data: TurbLog10 by Time  
Bartlett's K-squared = 0.017795, df = 1, p-value = 0.8939
```

The Bartlett's test indicates that the variances are equal. We can now perform the *t*-test on the transformed data.

Test statistic

```
t.test(TurbLog10 ~ Time, data = turbidity, var.equal = TRUE)
```

Two Sample t-test

```
data: TurbLog10 by Time  
t = -3.7169, df = 17, p-value = 0.001714  
alternative hypothesis: true difference in means between group After and group  
Before is not equal to 0  
95 percent confidence interval:  
-1.302095 -0.359140  
sample estimates:  
mean in group After mean in group Before  
0.683656 1.514274
```

P-value

The p-value is 0.007, which is less than 0.05, so we reject the null hypothesis that the mean turbidity values before and after the new filtering process was installed are equal.

Conclusion

We want to make a biological conclusion based on the results. To do this, we need to back-transform the mean difference and the confidence intervals so that the units are in NTU, not \log_{10} NTU.

Mean difference

```
10^(1.514274 - 0.683656)
```

```
[1] 6.770457
```

Lower and upper CI


```
10^(0.359140)
```

```
[1] 2.286336
```

```
10^(1.302095)
```

```
[1] 20.04911
```

The mean turbidity values before and after the new filtering process was installed are significantly different ($t = -3.7$, $df = 17$, $p = 0.007$). The new filtering process has effectively reduced turbidity by 7 NTU on average. We are 95% confident that the new filtering process will reduce water turbidity by between 2.3 and 20 NTU.

Bonus take home exercises

Use the HATPC framework to answer the following questions

1. Hypothesis
2. Assumptions
3. Test (statistic)
4. P-value
5. Conclusion

Exercise 1: Tooth growth

Tooth GRowth is an inbuilt dataset that shows the effect of vitamin c in Guinea pig tooth growth. It has three variables:

- len = tooth length
- supp = type of supplement (Orange juice or ascorbic acid)
- dose = mg/day given to the guinea pigs

```
head(ToothGrowth)
```

```
  len supp dose
1  4.2   VC  0.5
2 11.5   VC  0.5
3  7.3   VC  0.5
4  5.8   VC  0.5
5  6.4   VC  0.5
6 10.0   VC  0.5
```

```
str(ToothGrowth)
```

```
'data.frame':  60 obs. of  3 variables:
 $ len : num  4.2 11.5 7.3 5.8 6.4 10 11.2 11.2 5.2 7 ...
 $ supp: Factor w/ 2 levels "OJ","VC": 2 2 2 2 2 2 2 2 2 2 ...
 $ dose: num  0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 ...
```

Using the HATPC framework, test whether the type of supplement affects tooth length.

Solution

1. Hypothesis


$$H_0: \mu_{OJ} = \mu_{VC}$$

$$H_1: \mu_{OJ} \neq \mu_{VC}$$

2. Assumptions

Normality


```
library(tidyverse)
# qqplot
ggplot(ToothGrowth, aes(sample = len)) +
  stat_qq() +
  stat_qq_line() +
  facet_wrap(~supp)
```



The data appear to be normally distributed, as most of the dots follow the line

Homogeneity of variance

```
#Boxplots  
ggplot(ToothGrowth, aes(x = supp, y = len))+  
  geom_boxplot()
```



```
# Bartlett test
bartlett.test(len~supp, data = ToothGrowth)
```

Bartlett test of homogeneity of variances

data: len by supp
Bartlett's K-squared = 1.4217, df = 1, p-value = 0.2331

The two groups have equal variance

3. Test statistic

```
t.test(len~supp, var.equal = TRUE, data = ToothGrowth)
```

Two Sample t-test

data: len by supp
t = 1.9153, df = 58, p-value = 0.06039
alternative hypothesis: true difference in means between group OJ and group VC
is not equal to 0
95 percent confidence interval:

```

-0.1670064  7.5670064
sample estimates:
mean in group OJ mean in group VC
      20.66333      16.96333

```

4. P-value

Since the p-value > 0.05, we retain the null hypothesis

5. Conclusion

The mean tooth length was the same for each treatment group. The type of supplement does not have a significant effect on guinea pig tooth length

Exercise 2: Adelie penguin bill length

For this exercise, we will be using a subset of the palmer penguins dataset.

```

library(palmerpenguins)
library(tidyverse)

adelie <- penguins%>%
  filter(species == "Adelie")%>%
  na.omit()%>%
  droplevels()

head(adelie)

```

```

# A tibble: 6 × 8
  species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
  <fct>   <fct>         <dbl>         <dbl>             <int>         <int>
1 Adelie Torgersen      39.1           18.7             181          3750
2 Adelie Torgersen      39.5           17.4             186          3800
3 Adelie Torgersen      40.3           18              195          3250
4 Adelie Torgersen      36.7           19.3             193          3450
5 Adelie Torgersen      39.3           20.6             190          3650
6 Adelie Torgersen      38.9           17.8             181          3625
# i 2 more variables: sex <fct>, year <int>

```

Using the HATPC framework, test whether male and female penguins have the same length bill

Solution

1. Hypothesis

$$H_0: \mu_{fbilllength} = \mu_{mbilllength}$$


$$H_1: \mu_{fbilllength} \neq \mu_{mbilllength}$$

2. Assumptions

Normality

```
library(tidyverse)

# qqplot
ggplot(adelie, aes(sample = bill_length_mm))+
  geom_qq()+
  geom_qq_line()+
  facet_wrap(~sex)
```



```
#shapiro.test
tapply(adelie$bill_length_mm, adelie$sex, shapiro.test)
```

```
$female
```

Shapiro-Wilk normality test

```
data: X[[i]]
```

```
W = 0.99117, p-value = 0.8952
```

```
$male
```

```
Shapiro-Wilk normality test
```


```
data: X[[i]]
```

```
W = 0.98613, p-value = 0.6067
```

The two groups appear to be normally distributed, both according to the qqplot and the shapiro test

Homogeneity of variance

```
#boxplot  
ggplot(adelie, aes(x = sex, y = bill_length_mm))+  
  geom_boxplot()
```



```
#bartlett test  
bartlett.test(bill_length_mm~sex, data = adelie)
```

Bartlett test of homogeneity of variances

```
data: bill_length_mm by sex
Bartlett's K-squared = 0.95062, df = 1, p-value = 0.3296
```

The two groups have equal variance

3. Test (statistic)

```
# t.test
t.test(bill_length_mm~sex, var.equal = TRUE, data = adelie)
```

Two Sample t-test

```
data: bill_length_mm by sex
t = -8.7765, df = 144, p-value = 4.44e-15
alternative hypothesis: true difference in means between group female and group
male is not equal to 0
95 percent confidence interval:
-3.838435 -2.427319
sample estimates:
mean in group female    mean in group male
      37.25753           40.39041
```

4. P-value

Since $p < 0.05$, we reject the null hypothesis

5. Conclusion

Male and female adelie penguins have different mean bill lengths. Males have longer beaks than females

Exercise 3: Penguin body mass

For this exercise, we will use a subset of the palmer penguins data set again. This time, we will be comparing two different penguin species.

```
library(palmerpenguins)

penguins2<- penguins%>%
  filter(species != "Adelie")%>%
  na.omit()%>%
  droplevels()
```


Using the HATPC framework, test whether chinstrap and gentoo penguins have different body masses.

Solution

1. Hypothesis

$$H_0: \mu_{chinstrap} = \mu_{gentoo}$$

$$H_1: \mu_{chinstrap} \neq \mu_{gentoo}$$


2. Assumptions

Normality

```
library(tidyverse)

#qqplot

ggplot(penguins2, aes(sample = body_mass_g))+
  geom_qq()+
  geom_qq_line()+
  facet_wrap(~species)
```



```
#shapiro test
```

```
tapply(penguins2$body_mass_g, penguins2$species, shapiro.test)
```

```
$Chinstrap
```

```
Shapiro-Wilk normality test
```

```
data: X[[i]]
```

```
W = 0.98449, p-value = 0.5605
```

```
$Gentoo
```

```
Shapiro-Wilk normality test
```

```
data: X[[i]]
```

```
W = 0.98606, p-value = 0.2605
```

The data appear to be normally distributed. The points in the qq plot follow the line pretty closely, and the shapiro tests also indicate a normal distribution

Homogeneity of variance

```
# bartletttest
```

```
bartlett.test(penguins2$body_mass_g~penguins2$species)
```

```
Bartlett test of homogeneity of variances
```

```
data: penguins2$body_mass_g by penguins2$species
```

```
Bartlett's K-squared = 5.6724, df = 1, p-value = 0.01723
```

According to the Bartlett test, the two samples do not have equal variances

3. Test (statistic)

```
t.test(body_mass_g~species, var.equal = FALSE, data = penguins2)
```

```
Welch Two Sample t-test
```

```
data: body_mass_g by species
```

```
t = -20.765, df = 169.62, p-value < 2.2e-16
```

```
alternative hypothesis: true difference in means between group Chinstrap and
group Gentoo is not equal to 0
95 percent confidence interval:
-1488.578 -1230.120
sample estimates:
mean in group Chinstrap    mean in group Gentoo
      3733.088              5092.437
```

4. P-value

Since $p > 0.05$, we reject the null hypothesis

5. Conclusion

Chinstrap and gentoo penguins have different mean body masses. Gentoo penguins are heavier than chinstrap penguins on average.